

****Introduction to Ray Serve****

Ray Serve allows for efficient multi-model serving by combining the benefits of microservices and monoliths.

****Model Composition****

- * Efficient Hardware resource sharing between models.
- * Independent scaling of models, allowing for flexible resource allocation.
- * Simplification of application testing and monitoring.
- * Example: Combining image pre-processing, classification, and detection models.

****Multi-Application****

- * Supports multiple independent applications on a single cluster.
- * Facilitates collaboration between teams and ensures independent upgrade cycles.
- * Example: Managing autonomous driving algorithms for different scenarios and environmental conditions.

****Multiplexing****

- * Addresses the challenge of serving a large number of models with limited Hardware resources.
- * Dynamically manages model loading and routing traffic to specific replicas with cached models.
- * Improves model cache hit rate, reducing latency and boosting throughput.
- * Example: Support for numerous language models in an inference platform.

****Case Studies****

- * **Samsara:** Reduced ML infrastructure costs by 50% using Ray Serve model composition.
- * **Unscale Endpoints:** Boosted throughput by 30% using Ray Serve multiplexing.
- * **Clary:** Enhanced model training speed by 80% and serving latency by 80% using Ray Serve.